

CONTEXT

INTRODUCTION

WORKING

DIAGRAM

DESCRIPTION

IMPLEMENTATION

CODE

BOUNDARY CONDITIONS

TEST CASES

GITHUB REPOSITORY

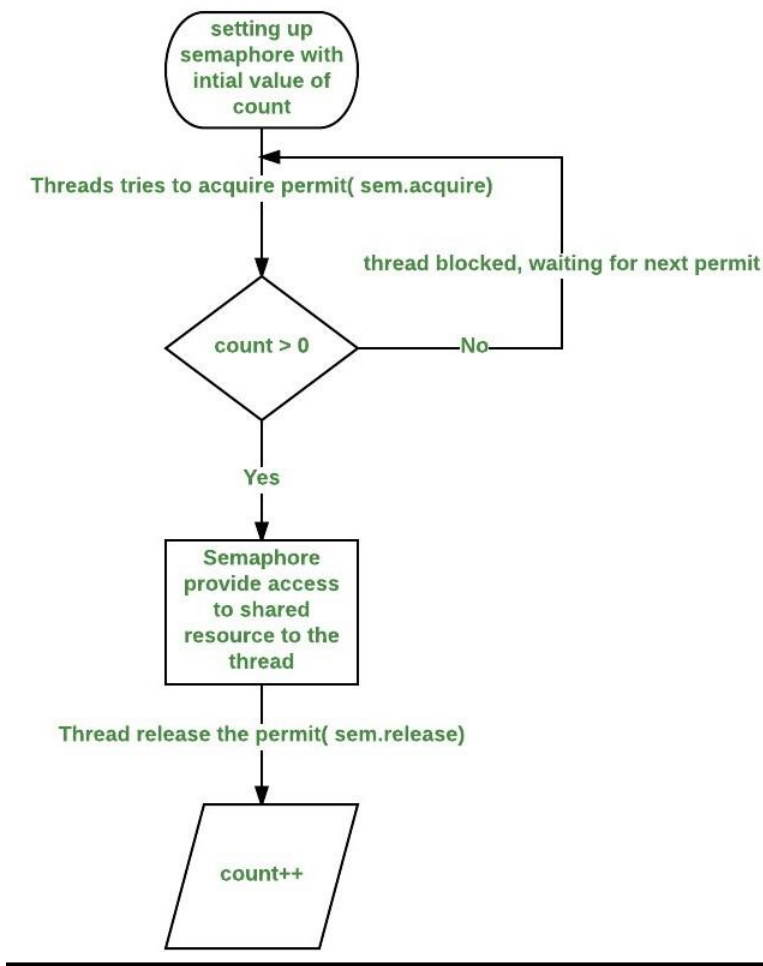
CONCLUSION

INTRODUCTION:

A semaphore controls access to a shared resource through the use of a counter. If the counter is greater than zero, then access is allowed. If it is zero, then access is denied. What the counter is counting are permits that allow access to the shared resource. Thus, to access the resource, a thread must be granted a permit from the semaphore.

WORKING:

In general, to use a semaphore, the thread that wants access to the shared resource tries to acquire a permit.



- If the semaphore's count is greater than zero, then the thread acquires a permit, which causes the semaphore's count to be decremented.
- Otherwise, the thread will be blocked until a permit can be acquired.
- When the thread no longer needs an access to the shared resource, it releases the permit, which causes the semaphore's count to be incremented.
- If there is another thread waiting for a permit, then that thread will acquire a permit at that time.

DESCRIPTION:

Semaphores in Process Synchronization

Semaphore was proposed by Dijkstra in 1965 which is a very significant technique to manage concurrent processes by using a simple integer value, which is known as a semaphore. Semaphore is simply a variable which is non-negative and shared between threads. This variable is used to solve the critical section problem and to achieve process synchronization in the multiprocessing environment.

Semaphores are of two types:

1. **Binary Semaphore** – This is also known as mutex lock. It can have only two values – 0 and 1. Its value is initialized to 1. It is used to implement the solution of critical section problem with multiple processes.
2. **Counting Semaphore** – Its value can range over an unrestricted domain. It is used to control access to a resource that has multiple instances.

Some point regarding P and V operation

1. P operation is also called wait, sleep or down operation and V operation is also called signal, wake-up or up operation.
2. Both operations are atomic and semaphore(s) is always initialized to one. Here atomic means that variable on which read, modify and update happens at the same time/moment with no pre-emption i.e. in between read, modify and update no other operation is performed that may change the variable.
3. A critical section is surrounded by both operations to implement process synchronization. See below image. critical section of Process P is in between P and V operation.

ADVANTAGES:

- Semaphores allow only one process into the critical section. They follow the mutual exclusion principle strictly and are much more efficient than some other methods of synchronization.
- There is no resource wastage because of busy waiting in semaphores as processor time is not wasted unnecessarily to check if a condition is fulfilled to allow a process to access the critical section.
- Semaphores are implemented in the machine independent code of the microkernel. So they are machine independent.

DISADVANTAGES:

- Semaphores are complicated so the wait and signal operations must be implemented in the correct order to prevent deadlocks.
- Semaphores are impractical for large scale use as their use leads to loss of modularity. This happens because the wait and signal operations prevent the creation of a structured layout for the system.
- Semaphores may lead to a priority inversion where low priority processes may access the critical section first and high priority processes later.

IMPLEMENTATION:

Implementation of binary semaphores:

```
struct semaphore {
    enum value(0, 1);

    // q contains all Process Control Blocks (PCBs)
    // corresponding to processes got blocked
    // while performing down operation.
    Queue<process> q;
}

P(semaphore s)
{
    if (s.value == 1) {
        s.value = 0;
    }
    else {
        // add the process to the waiting queue
        q.push(P)
        sleep();
    }
}

V(Semaphore s)
{
    if (s.q is empty) {
        s.value = 1;
    }
    else {
        // select a process from waiting queue
        q.pop();
        wakeup();
    }
}
```

The description above is for binary semaphore which can take only two values 0 and 1 and ensure the mutual exclusion. There is one other type of semaphore called counting semaphore which can take values greater than one.

Limitations

1. One of the biggest limitation of semaphore is priority inversion.
2. Deadlock, suppose a process is trying to wake up another process which is not in sleep state. Therefore a deadlock may block indefinitely.
3. The operating system has to keep track of all calls to wait and to signal the semaphore.

Problem in this implementation of semaphore

Whenever any process waits then it continuously checks for semaphore value (look at this line while (s==0); in P operation) and waste CPU cycle. To avoid this another implementation is provided below.

Implementation of counting semaphore

```
struct Semaphore {
    int value;

    // q contains all Process Control Blocks(PCBs)
    // corresponding to processes got blocked
    // while performing down operation.
    Queue<process> q;
}

P(Semaphore s)
{
    s.value = s.value - 1;
    if (s.value < 0) {

        // add process to queue
        // here p is a process which is currently executing
        q.push(p);
        block();
    }
    else
        return;
}

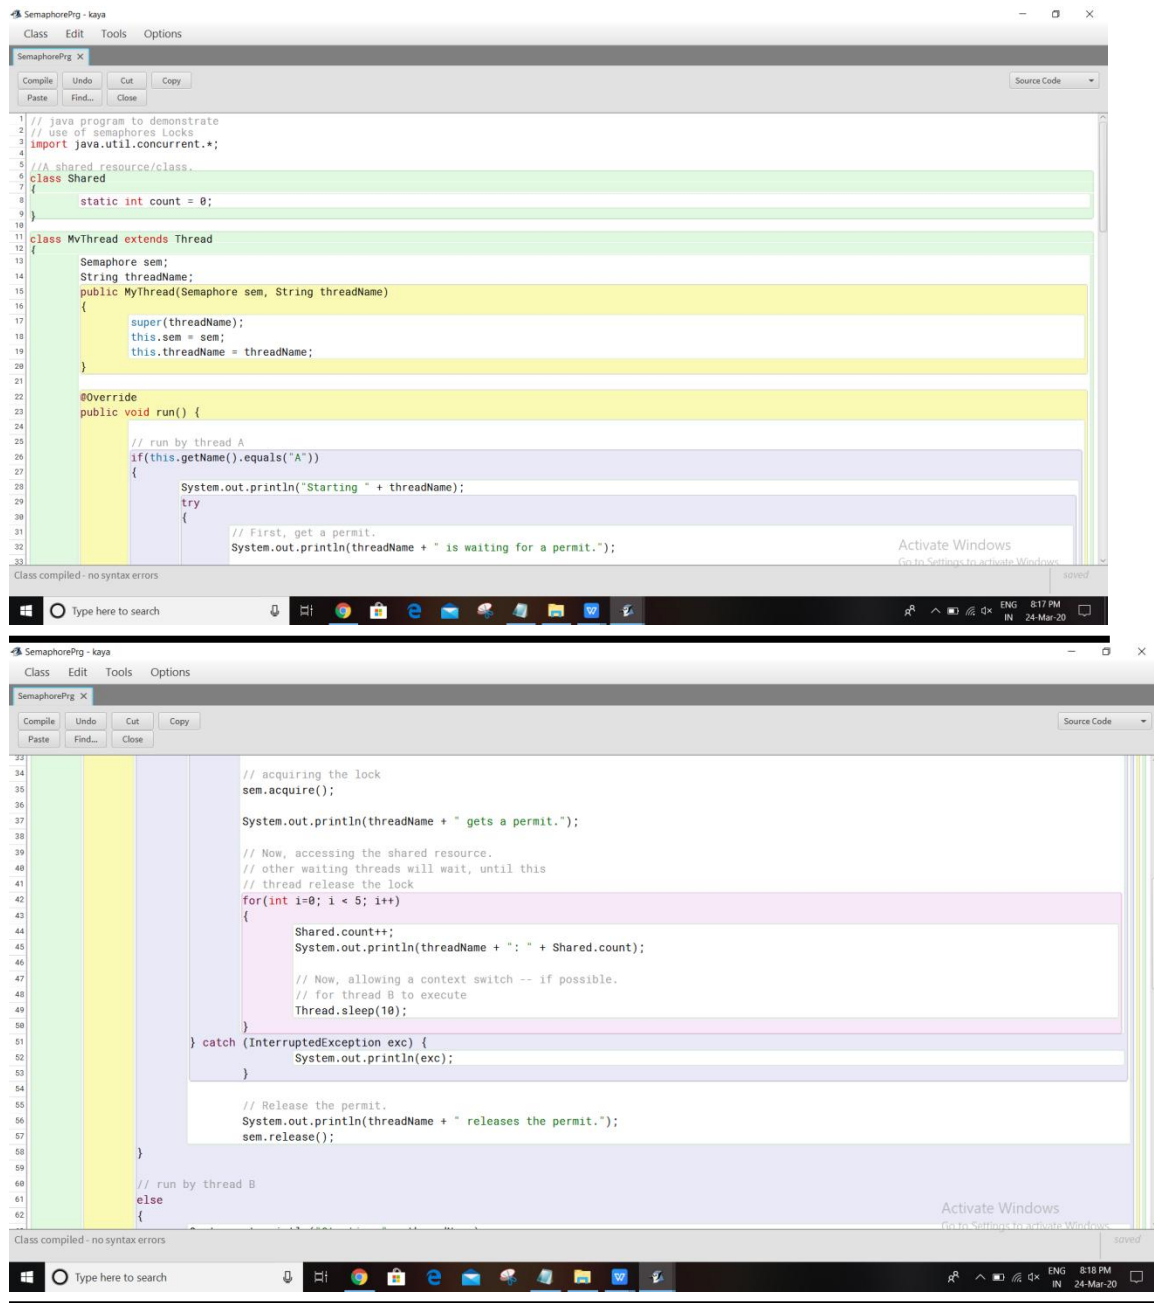
V(Semaphore s)
{
    s.value = s.value + 1;
    if (s.value <= 0) {

        // remove process p from queue
        q.pop();
    }
}
```

```
        wakeup(p);  
    }  
    else  
        return;  
}
```

In this implementation whenever process waits it is added to a waiting queue of processes associated with that semaphore. This is done through system call `block()` on that process. When a process is completed it calls `signal` function and one process in the queue is resumed. It uses `wakeup()` system call.

CODE:



```
1 // java program to demonstrate
2 // use of semaphores locks
3 import java.util.concurrent.*;
4
5 //A shared resource/class.
6 class Shared
7 {
8     static int count = 0;
9 }
10
11 class MyThread extends Thread
12 {
13     Semaphore sem;
14     String threadName;
15     public MyThread(Semaphore sem, String threadName)
16     {
17         super(threadName);
18         this.sem = sem;
19         this.threadName = threadName;
20     }
21
22     @Override
23     public void run() {
24
25         // run by thread A
26         if(this.getName().equals("A"))
27         {
28             System.out.println("Starting " + threadName);
29             try
30             {
31                 // First, get a permit.
32                 System.out.println(threadName + " is waiting for a permit.");
33
34                 // acquiring the lock
35                 sem.acquire();
36
37                 System.out.println(threadName + " gets a permit.");
38
39                 // Now, accessing the shared resource.
40                 // other waiting threads will wait, until this
41                 // thread release the lock
42                 for(int i=0; i < 5; i++)
43                 {
44                     Shared.count++;
45                     System.out.println(threadName + " : " + Shared.count);
46
47                     // Now, allowing a context switch -- if possible.
48                     // for thread B to execute
49                     Thread.sleep(10);
50                 }
51             } catch (InterruptedException exc) {
52                 System.out.println(exc);
53             }
54
55             // Release the permit.
56             System.out.println(threadName + " releases the permit.");
57             sem.release();
58         }
59
60         // run by thread B
61         else
62         {
63             // ... (code for thread B) ...
64         }
65     }
66 }
```

SemaphorePrg - kaya

Class Edit Tools Options

SemaphorePrg X

Compile Undo Cut Copy
Paste Find... Close Source Code

```
63 System.out.println("Starting " + threadName);
64 try
65 {
66     // First, get a permit.
67     System.out.println(threadName + " is waiting for a permit.");
68
69     // acquiring the lock
70     sem.acquire();
71
72     System.out.println(threadName + " gets a permit.");
73
74     // Now, accessing the shared resource.
75     // other waiting threads will wait, until this
76     // thread release the lock
77     for(int i=0; i < 5; i++)
78     {
79         Shared.count--;
80         System.out.println(threadName + ": " + Shared.count);
81
82         // Now, allowing a context switch -- if possible.
83         // for thread A to execute
84         Thread.sleep(10);
85     }
86 } catch (InterruptedException exc) {
87     System.out.println(exc);
88 }
89 // Release the permit.
90 System.out.println(threadName + " releases the permit.");
91 sem.release();
92 }
```

Class compiled - no syntax errors

Activate Windows
Go to Settings to activate Windows

saved

Type here to search

ENG 5:18 PM
IN 24-Mar-20

SemaphorePrg - kaya

Class Edit Tools Options

SemaphorePrg X

Compile Undo Cut Copy
Paste Find... Close Source Code

```
95 // Driver class
96 public class SemaphorePrg
97 {
98     public static void main(String args[]) throws InterruptedException
99     {
100         // creating a Semaphore object
101         // with number of permits 1
102         Semaphore sem = new Semaphore(1);
103
104         // creating two threads with name A and B
105         // Note that thread A will increment the count
106         // and thread B will decrement the count
107         MyThread mt1 = new MyThread(sem, "A");
108         MyThread mt2 = new MyThread(sem, "B");
109
110         // stating threads A and B
111         mt1.start();
112         mt2.start();
113
114         // waiting for threads A and B
115         mt1.join();
116         mt2.join();
117
118         // count will always remain 0 after
119         // both threads will complete their execution
120         System.out.println("count: " + Shared.count);
121     }
122 }
123
124 }
```

Class compiled - no syntax errors

Activate Windows
Go to Settings to activate Windows

saved

Type here to search

ENG 5:18 PM
IN 24-Mar-20

REPOSITORY:

CONCLUSION:

